

# **Chapter 4**

## **Threads & Concurrency**

Yunmin Go

School of CSEE



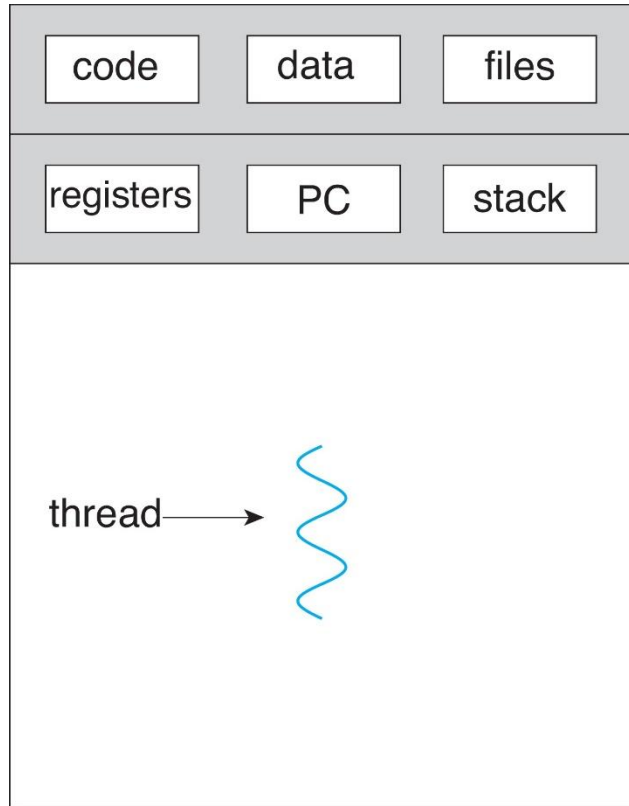
# Agenda

- Overview
- Parallelism and Concurrency
- Multithreading Models
- Thread Creation
- Thread Cancellation & Some Issues
- Implicit Threading

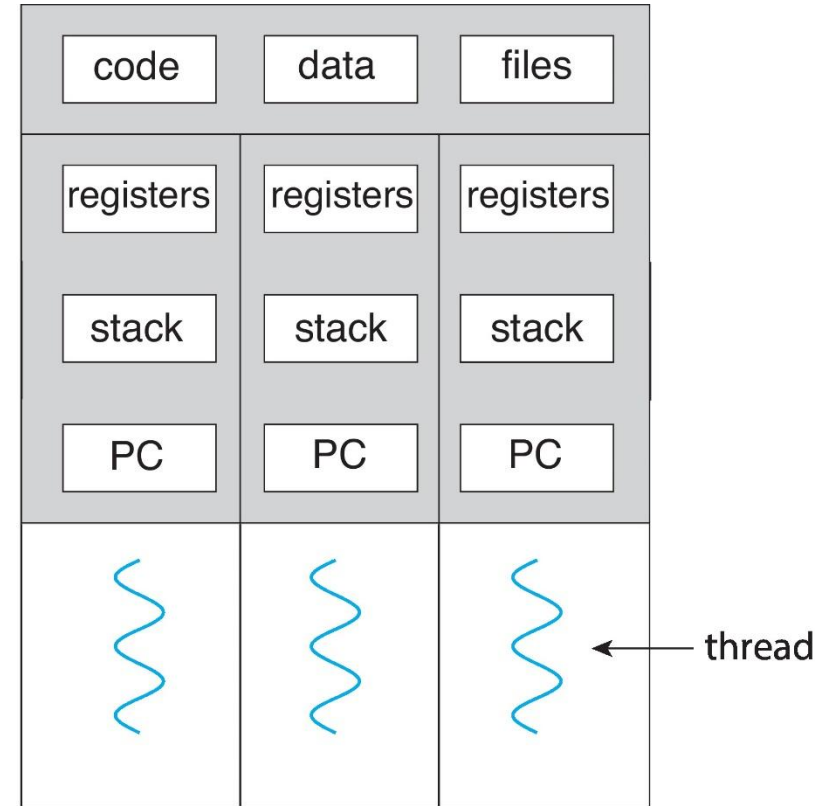
# Overview

- **Process**: program in execution
  - Each process occupies resources required for execution
- **Thread**: a way for a program to split itself into two or more simultaneously running tasks
  - Basic unit of CPU utilization
  - Smaller unit than process
  - Threads in a process share resources
- A thread is comprised of
  - Thread ID, program counter, register set, stack, etc.

# Single and Multithreaded Processes



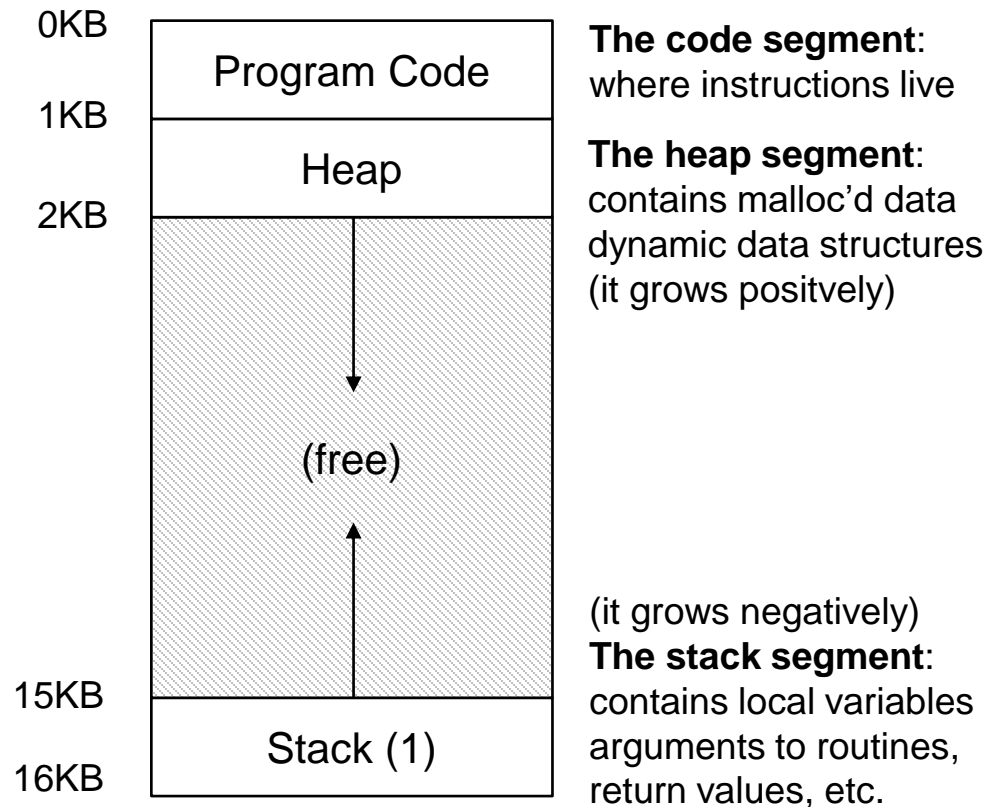
single-threaded process




multithreaded process

# Address Spaces

- There will be one stack per thread



<A Single-Threaded Address Space>



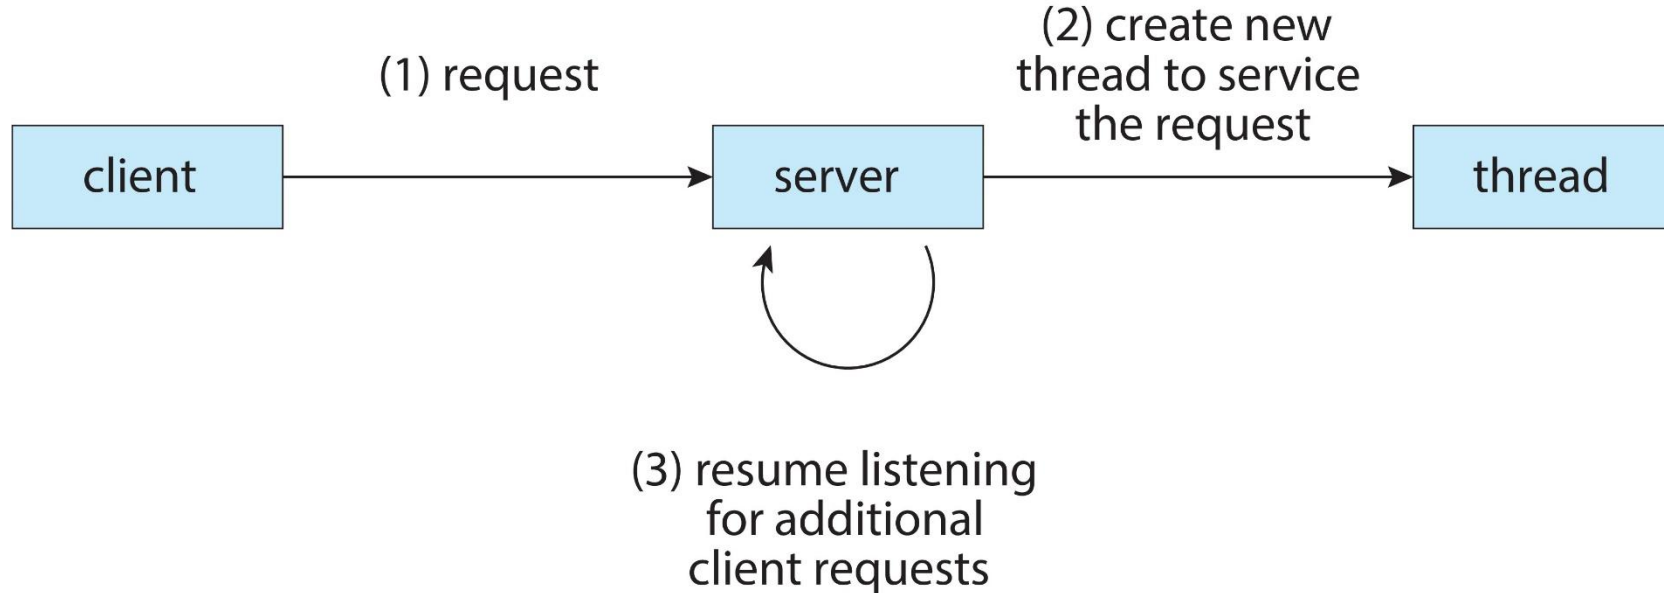
<Two threaded Address Space>

# Thread Control Block (TCB)

- **Thread Control Block (TCB)** is a data structure in the operating system kernel which contains thread-specific information needed to manage it.
- Examples of information in TCB (`thread_struct` in Linux)
  - Thread id
  - State of the thread (running, ready, waiting, start, done)
  - Stack pointer
  - Program counter
  - Thread's register values
  - Pointer to the process control block (PCB)

# Why Use Threads?

- Process creation is expensive in time and resource
  - ex) Web server accepting thousands of requests



# Why Use Threads?

## ■ Parallelism

- Single-threaded program: the task is straightforward, but slow.
- Multi-threaded program: natural and typical way to make programs run faster on modern hardware.
- **Parallelization:** The task of transforming standard single-threaded program into a program that does this sort of work on multiple CPUs.

- To avoid blocking program progress due to slow I/O.
  - Threading enables overlap of I/O with other activities within a single program.
  - It is much like multiprogramming did for processes across programs.



# Benefits

- **Responsiveness:** may allow continued execution if part of process is blocked, especially important for user interfaces
- **Resource Sharing:** threads share resources of process, easier than shared memory or message passing
- **Economy:** cheaper than process creation, thread switching lower overhead than context switching
- **Scalability:** process can take advantage of multiprocessor architectures

# Agenda

- Overview
- **Parallelism and Concurrency**
- Multithreading Models
- Thread Creation
- Thread Cancellation & Some Issues
- Implicit Threading


# Parallelism and Concurrency

- **Multicore** or **multiprocessor** systems putting pressure on programmers, challenges include:
  - Dividing activities
  - Balance
  - Data splitting
  - Data dependency
  - Testing and debugging
- **Parallelism** implies a system can perform more than one task simultaneously
- **Concurrency** supports more than one task making progress
  - Single processor / core, scheduler providing concurrency


# Parallelism and Concurrency

- Concurrency vs. Parallelism

- Concurrent execution on single-core system:




- Parallelism on a multi-core system:




# Types of Parallelism

- **Data parallelism**: distributes subsets of the same data across multiple cores, same operation on each



- **Task parallelism**: distributing threads across cores, each thread performing unique operation




# Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components

- S is serial portion, N processing cores

$$speedup \leq \frac{1}{S + \frac{(1 - S)}{N}}$$

- ex) Application is 75% parallel (25% serial), moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches 1 / S
- Serial portion of an application has disproportionate effect on performance gained by adding additional cores.




But does the law take into account contemporary multicore systems?

# Agenda

- Overview
- Parallelism and Concurrency
- **Multithreading Models**
- Thread Creation
- Thread Cancellation & Some Issues
- Implicit Threading

# User Threads and Kernel Threads

- **User thread:** thread supported by **thread library in user level**
  - Created by library function call (not system call)
  - Kernel is not concerned in user thread
  - Switching of user thread is faster than kernel thread.
- **Kernel thread:** thread supported by **kernel**
  - Created and managed by kernel
  - Scheduled by kernel
  - Cheaper than process
  - More expensive than user thread





# Kernel Thread

- Most operating system kernels are also typically multithreaded
  - Each thread performs a specific task, such as managing devices, memory management, or interrupt handling


```
mcnl@mcnl:~$ ps -ef
UID          PID    PPID  C   STIME TTY          TIME CMD
root           1        0  0   2021 ?        00:06:14 /lib/systemd/systemd --system --deserialize 37
root           2        0  0   2021 ?        00:00:03 [kthreadd]
root           3         2  0   2021 ?        00:00:00 [rcu_gp]
root           4         2  0   2021 ?        00:00:00 [rcu_par_gp]
root           6         2  0   2021 ?        00:00:00 [kworker/0:0H-kblockd]
root           9         2  0   2021 ?        00:00:00 [mm_percpu_wq]
root          10         2  0   2021 ?        00:00:11 [ksoftirqd/0]
root          11         2  0   2021 ?        00:11:21 [rcu_sched]
root          12         2  0   2021 ?        00:00:23 [migration/0]
root          13         2  0   2021 ?        00:00:00 [idle_inject/0]
root          14         2  0   2021 ?        00:00:00 [cpuhp/0]
root          15         2  0   2021 ?        00:00:00 [cpuhp/1]
root          16         2  0   2021 ?        00:00:00 [idle_inject/1]
root          17         2  0   2021 ?        00:00:24 [migration/1]
root          18         2  0   2021 ?        00:00:02 [ksoftirqd/1]
root          20         2  0   2021 ?        00:00:00 [kworker/1:0H-kblockd]
root          21         2  0   2021 ?        00:00:00 [cpuhp/2]
root          22         2  0   2021 ?        00:00:00 [idle_inject/2]
root          23         2  0   2021 ?        00:00:26 [migration/2]
root          24         2  0   2021 ?        00:00:01 [ksoftirqd/2]
root          26         2  0   2021 ?        00:00:00 [kworker/2:0H-kblockd]
root          27         2  0   2021 ?        00:00:00 [cpuhp/3]
root          28         2  0   2021 ?        00:00:00 [idle_inject/3]
root          29         2  0   2021 ?        00:00:25 [migration/3]
root          30         2  0   2021 ?        00:00:01 [ksoftirqd/3]
```

# Multithreading Models

- Major issue: correspondence between user threads and kernel threads
- Multithreading Models
  - Many-to-One
  - One-to-One
  - Many-to-Many (Two-level)


# Many-to-One

- Many user-level threads mapped to single kernel thread
- One thread blocking causes all to block
- Multiple threads may not run in parallel on multicore system because only one may be in kernel at a time
- Few systems currently use this model
- Examples:
  - Solaris Green Threads
  - Early version of Java




# One-to-One

- Each user-level thread maps to kernel thread
- Creating a user-level thread creates a kernel thread
- More concurrency than many-to-one
- Number of threads per process sometimes restricted due to overhead
- Examples
  - Windows
  - Linux




# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise not very common



Many-to-Many Model



Two-level Model

# Agenda

- Overview
- Parallelism and Concurrency
- Multithreading Models
- **Thread Creation**
- Thread Cancellation & Some Issues
- Implicit Threading

# Thread Libraries

- **Thread library** provides programmer with API for creating and managing threads
- Two primary ways of implementing
  - User level library (entirely in user space with no kernel support)
  - Kernel-level library supported by the OS

# Pthreads

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
  - ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)



# Pthread Create

- `pthread_create()`: Create a new thread

```
#include <pthread.h>
int pthread_create(pthread_t *thread, const pthread_attr_t *attr,
                  void *(*start_routine) (void *), void *arg);
```

- Starts a new thread in the calling process
- *thread*: new thread ID on success
- *attr*: to set thread attribute (stack size, scheduling priority...). NULL for default value.
- *start\_routine*: the routine that the thread will execute when it is created
- *arg*: argument passed to *start\_routine*
- Return value
  - Success: 0
  - Error: error number

# Pthread Join

- `pthread_join()`: Join with a terminated thread

```
#include <pthread.h>
int pthread_join(pthread_t thread, void **retval);
```

- It blocks the calling thread until the specified thread returns
- When a parent thread returns, all the child thread are terminated
- *thread*: specified thread to terminate
- *retval*: return value from the terminated thread (a pointer returned by reference)
- Return value
  - Success: 0
  - Error: error number

# Compiling and Running

- To compile them, you must include the header `pthread.h`
  - Explicitly link with the pthreads library, by adding the `-pthread` flag.

```
$ gcc -o main main.c -Wall -pthread
```

- The recent gcc automatically links with the pthread library.

# Example: Thread Creation

```
#include <stdio.h>
#include <pthread.h>

void *mythread (void *arg) {
    printf ("%s\n", (char *) arg);
    return NULL;
}

int main (int argc, char *argv[]) {
    pthread_t p1, p2;
    int rc;
    printf("main: begin\n");
    pthread_create(&p1, NULL, mythread, "A");
    pthread_create(&p2, NULL, mythread, "B");

    // join waits for the threads to finish
    pthread_join(p1, NULL);
    pthread_join(p2, NULL);
    printf("main: end\n");
    return 0;
}
```

thread\_create.c

```
yunmin@yunmin:~/ch4$ ./thread_create
main: begin
A
B
main: end
```

## \* Trace #1

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
waits for T1	runs prints "A" returns	
		runs prints "B" returns
waits for T2		
prints "main: end"		

## \* Trace #2

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
		runs prints "A" returns
		runs prints "B" returns
creates Thread 2		
waits for T1		
<i>returns immediately; T1 is done</i>		
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

## \* Trace #3

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1		
creates Thread 2		
		runs prints "B" returns
waits for T1		
	runs prints "A" returns	
waits for T2		
<i>returns immediately; T2 is done</i>		
prints "main: end"		

# Example: Passing Argument to Thread

```
#include <stdio.h>
#include <pthread.h>

void *thread_summation(void * arg);
int sum = 0;

void *thread_summation(void * arg) {
    int start = ((int*)arg)[0];
    int end = ((int*)arg)[1];
    while (start <= end)
    {
        sum+=start;
        start++;
    }
    return NULL;    // or pthread_exit(0);
}
```

thread\_sum.c

```
int main(int argc, char *argv[]) {
    pthread_t id_t1, id_t2; // thread id
    int range1[] = {1, 5};
    int range2[] = {6, 10};

    // create thread
    pthread_create(&id_t1, NULL, thread_summation, (void *)range1);
    pthread_create(&id_t2, NULL, thread_summation, (void *)range2);

    // wait for the thread to exit
    pthread_join(id_t1, NULL);
    pthread_join(id_t2, NULL);
    printf("result: %d \n", sum);
    return 0;
}
```

Result?

# Example: Return Value from Thread

```
thread_ret.c
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

typedef struct __myret_t {
    int x;
    int y;
} myret_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    myret_t *r = malloc(sizeof(myret_t));
    r->x = 1;
    r->y = 2;
    return (void *) r;
}
```

```
int main(int argc, char *argv[]) {
    int rc;
    pthread_t p;
    myret_t *m;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    pthread_create(&p, NULL, mythread, &args);
    pthread_join(p, (void **) &m); // this thread has been
                                    // waiting inside of the
                                    // pthread_join() routine.

    printf("returned %d %d\n", m->x, m->y);
    return 0;
}
```

Result?

# Dangerous Code

- Be careful with how values are returned from a thread

```
void *mythread(void *arg) {  
    myarg_t *m = (myarg_t *) arg;  
    printf("%d %d\n", m->a, m->b);  
    myret_t r; // ALLOCATED ON STACK: BAD!  
    r.x = 1;  
    r.y = 2;  
    return (void *) &r;  
}
```

- When the variable *r* returns, it is automatically de-allocated

# Multiple Threads

- Pthreads code for joining 10 threads

```
#define NUM_THREADS 10

// an array of threads to be joined upon
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_create(&workers[i], NULL, thread_worker, (void *)arg);

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



# Agenda

- Overview
- Parallelism and Concurrency
- Multithreading Models
- Thread Creation
- **Thread Cancellation & Some Issues**
- Implicit Threading

# Pthread Cancel

- `pthread_join()`: Send a cancellation request to a thread

```
#include <pthread.h>
int pthread_cancel(pthread_t thread);
```

- This function sends a cancellation request to the thread
- Whether and when the target thread reacts to the cancellation request depends on two attributes that are under the control of that thread: its cancelability state and type.
- Cancelability state is determined by `pthread_setcancelstate()`
- Cancellation type is determined by `pthread_setcanceltype()`
- *thread*: thread ID to cancel
- Return value
  - Success: 0
  - Error: nonzero error number

# Example: Pthread Cancel

```
thread_cancel.c

#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <pthread.h>

int counter = 0;          // Counter
pthread_t tmp_thread;     // Thread ID for func2

void* func1(void* arg) {
    while (1) {
        printf("Thread #1 (counter=%d)\n", counter);
        if (counter == 5) {
            // for cancel thread_two
            pthread_cancel(tmp_thread);
            // for exit from thread_one
            pthread_exit(NULL);
        }
        sleep(1); // sleep 1 second
    }
}

void* func2(void* arg) {
    // get thread ID
    tmp_thread = pthread_self();
    while (1) {
        printf("Thread #2 (counter=%d)\n", counter);
        counter++;
        sleep(1); // sleep 1 second
    }
}
```

```
int main() {
    pthread_t thread_one, thread_two;

    // create thread_one and thread_two
    pthread_create(&thread_one, NULL, func1, NULL);
    pthread_create(&thread_two, NULL, func2, NULL);

    // waiting for when threads are completed
    pthread_join(thread_one, NULL);
    pthread_join(thread_two, NULL);

    return 0;
}
```

```
yunmin@yunmin:~/ch4$ gcc thread_cancel.c -o thread_cancel -lpthread
yunmin@yunmin:~/ch4$ ./thread_cancel
Thread #2 (counter=0)
Thread #1 (counter=1)
Thread #2 (counter=1)
Thread #1 (counter=2)
Thread #1 (counter=2)
Thread #2 (counter=2)
Thread #1 (counter=3)
Thread #2 (counter=3)
Thread #2 (counter=4)
Thread #1 (counter=5)
```

# Thread Cancellation

- Problem with thread cancellation
  - A thread share the resource with other threads  
cf. A process has its own resource.  
→ A thread can be cancelled while it updates data shared with other threads
- If thread has cancellation disabled, cancellation remains pending until thread enables it.
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

# Thread Cancellation

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

## ■ Setting thread cancellation type

- `pthread_t pthread_setcanceltype(int type, int *oldtype);`
  - *type*: PTHREAD\_CANCEL\_ASYNCHRONOUS (Asynchronous cancellation) or PTHREAD\_CANCEL\_DEFERRED (Deferred cancellation)
  - Default type is deferred
    - Cancellation only occurs when thread reaches **cancellation point** i.e. `pthread_testcancel()`, then cleanup handler is invoked
    - Safer than asynchronous cancellation

## ■ Enabling/disabling thread cancellation

- `pthread_t pthread_setcancelstate(int state, int *oldstate);`
  - *state*: PTHREAD\_CANCEL\_DISABLE or PTHREAD\_CANCEL\_ENABLE
  - PTHREAD\_CANCEL\_DISABLE: cancellation remains pending until thread enables it

# Example: Thread Cancellation

```
#include <pthread.h>
#include <stdio.h>
#include <unistd.h>
#include <limits.h>

unsigned int counter;

void* threadFunction(void* arg) {
    pthread_setcanceltype(*(int*)arg, NULL);
    counter = 0;
    while (1) {
        printf("Running... %u \n", counter);
        while (counter < 4000000000) {
            counter++;
            // if (counter % 500000000 == 0)
            //     pthread_testcancel();
        }
    }
    return NULL;
}
```

```
int main() {
    pthread_t thread;
    int async = PTHREAD_CANCEL_ASYNCHRONOUS;
    int deferred = PTHREAD_CANCEL_DEFERRED;
    printf("Maximum Value of Unsigned Long=%u \n", UINT_MAX);

    // Asynchronous cancellation
    pthread_create(&thread, NULL, threadFunction, &async);
    sleep(3);
    pthread_cancel(thread);
    pthread_join(thread, NULL);
    printf("Thread with asynchronous cancellation stopped. "
           "counter=%u \n", counter);

    // Deferred cancellation
    pthread_create(&thread, NULL, threadFunction, &deferred);
    sleep(3);
    pthread_cancel(thread);
    pthread_join(thread, NULL);
    printf("Thread with deferred cancellation stopped. "
           "counter=%u \n", counter);
    return 0;
}
```

```
yunmin@yunmin:~/ch4$ ./thread_setcancel
Maximum Value of Unsigned Long=4294967295
Running... 0
Thread with asynchronous cancellation stopped. counter=2585900876
Running... 0
Running... 4000000000
Thread with deferred cancellation stopped. counter=4000000000
```

# Thread-Local Storage

- In a process, all threads share global variables
- **Thread-local storage (TLS)** allows each thread to have its own copy of data

```
__thread int tls;           // on pthread
```

- Each thread has its own 'int tls' variable
- Different from local variables
  - Local variables visible only during single function invocation
  - TLS visible across function invocations
- Similar to static data
  - TLS is unique to each thread
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

# Thread-Local Storage in pthread

```
thread_tls.c

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <pthread.h>

#define THREADS 3

__thread int tls;
int global;

void *func(void *arg)
{
    int num = *((int*)arg);
    tls = num;
    global = num;
    sleep(1);
    printf("Thread = %d tls = %d global = %d\n",
        num, tls, global);
}
```

```
int main() {
    int ret;
    pthread_t thread[THREADS];
    int num;

    for (num = 0; num < THREADS; num++) {
        ret = pthread_create(&thread[num], NULL,
                            &func, (void*)&num);

        if (ret) {
            printf("error pthread_create\n");
            exit(1);
        }
    }

    for (num = 0; num < THREADS; num++) {
        ret = pthread_join(thread[num], NULL);
        if (ret) {
            printf("error pthread_join\n");
            exit(1);
        }
    }

    return 0;
}
```



# fork() and exec()

- fork() on multithreaded process
  - Duplicates all threads in the process?
  - Duplicates only corresponding thread?
  - UNIX supports two versions of fork
    - fork(), fork1()
- exec() on multithreaded process
  - Replace entire process including all threads

# Linux Threads

- Linux refers to them as ***tasks*** rather than ***threads***
- Thread creation is done through `clone()` system call
- `clone()` allows a child task to share the address space of the parent task (process)
  - Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- `struct task_struct` points to process data structures (shared or unique)

# Agenda


- Overview
- Parallelism and Concurrency
- Multithreading Models
- Thread Creation
- Thread Cancellation & Some Issues
- **Implicit Threading**

# Implicit Threading

- **Implicit Threading**: Creation and management of threads done by compilers and run-time libraries rather than programmers
- Three methods explored
  - Thread pools
  - Fork Join model
  - OpenMP
- Other methods include Grand Central Dispatch and Microsoft Threading Building Blocks (TBB), etc

# Thread Pools

- Create a number of threads in a pool where they await work
- Advantages
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool
  - Separating task to be performed from mechanics of creating task allows different strategies for running task
    - i.e. tasks could be scheduled to execute after a time delay or to execute periodically



# Java Thread Pools

- Supported by Java executor framework in `java.util.concurrent` package
- Three factory methods for creating thread pools in `Executors` class:
  - `static ExecutorService newSingleThreadExecutor()`
    - Creates a pool of size 1
  - `static ExecutorService newFixedThreadPool(int size)`
    - Creates a thread pool with a specified number of threads
  - `static ExecutorService newCachedThreadPool()`
    - Creates an unbounded thread pool, reusing in many instances

# OpenMP

- Provides support for parallel programming in shared-memory environments
  - Set of compiler directives and an API for C, C++, FORTRAN
  - Identifies **parallel regions**: blocks of code that can run in parallel
- Create as many threads as there are cores

```
#pragma omp parallel           // each runs the statement
    printf("Hello, World!\n");
```

- Run for loop in parallel

```
#pragma omp parallel for       // unroll loop over cores
    for(i=0; i<N; i++) {
        c[i] = a[i] + b[i];
    }
```

# Example: OpenMP

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */
    #pragma omp parallel
    {
        printf("I am a parallel region (thread=%d)\n", omp_get_thread_num());
    }

    /* sequential code */
    return 0;
}
```

Results (depends on your machine)

```
yunmin@mcn1-server:~/workspace/os/ch4$ gcc openmp.c -o openmp -fopenmp
yunmin@mcn1-server:~/workspace/os/ch4$ ./openmp
I am a parallel region (thread=81)
I am a parallel region (thread=18)
I am a parallel region (thread=85)
I am a parallel region (thread=54)
I am a parallel region (thread=15)
I am a parallel region (thread=41)
I am a parallel region (thread=11)
I am a parallel region (thread=94)
I am a parallel region (thread=50)
I am a parallel region (thread=58)
```