

Chapter 5

CPU Scheduling

Yunmin Go

School of CSEE



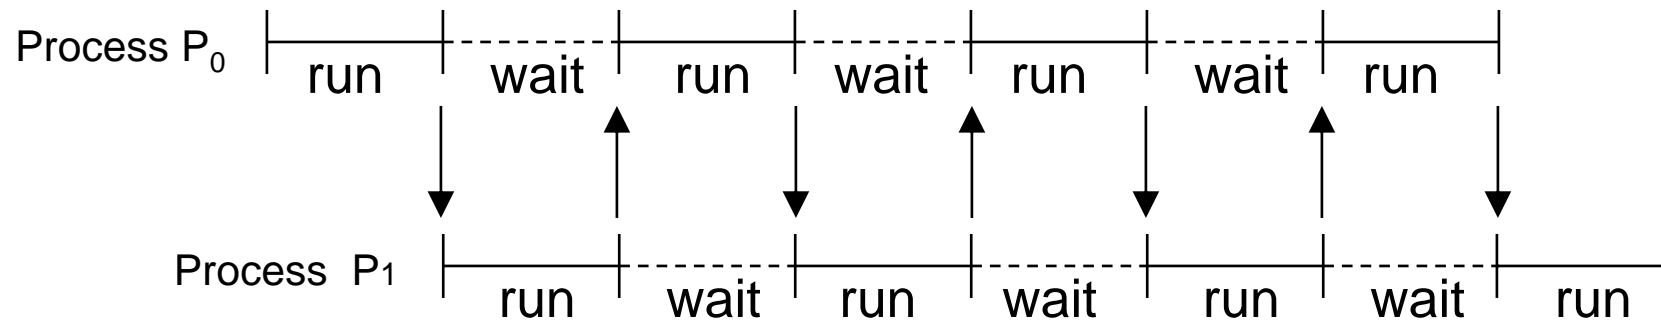
Agenda

- Basic Concepts
- Scheduling Algorithms
- Real-Time CPU Scheduling
- Linux Scheduler



Basic Concepts

- Motivation: maximum CPU utilization obtained with multiprogramming and multitasking
 - Resources (including CPU) are shared among processes



Multiprogramming

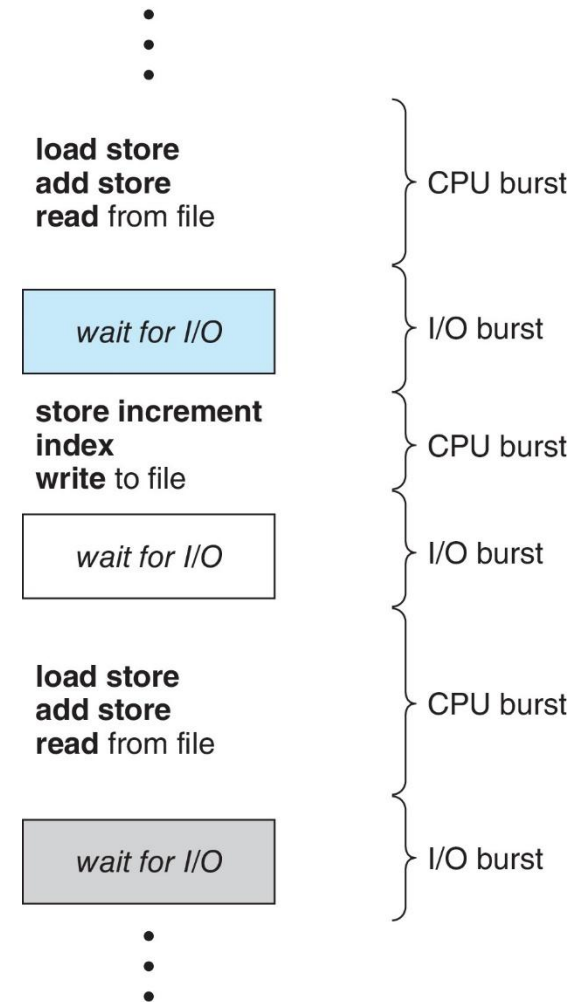
CPU-IO Burst Cycle

- Process execution consists of cycle of CPU execution and I/O wait
 - First and last bursts are CPU bursts

- Types of processes

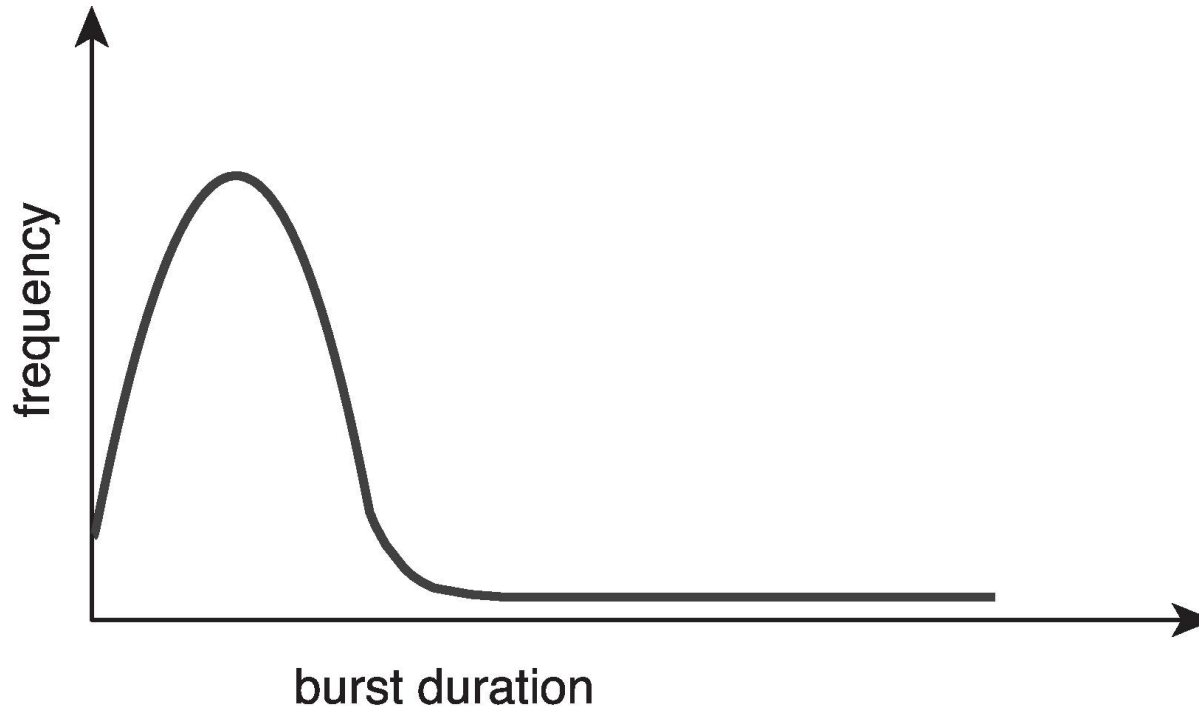
- I/O-bound process
 - Consists of many short CPU bursts
- CPU-bound process
 - Consists of a few long CPU bursts

CPU burst followed by **I/O burst**



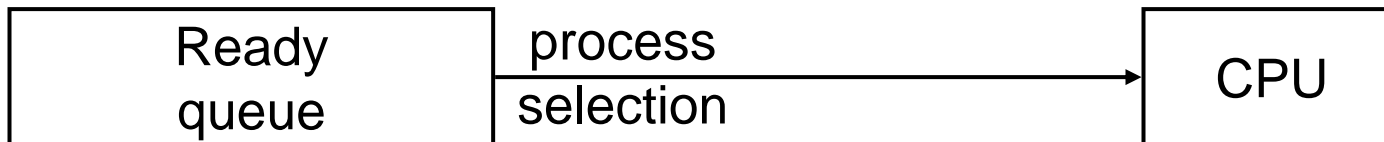
Histogram of CPU-burst Time

- CPU burst distribution is of main concern
 - Large number of short bursts
 - Small number of longer bursts



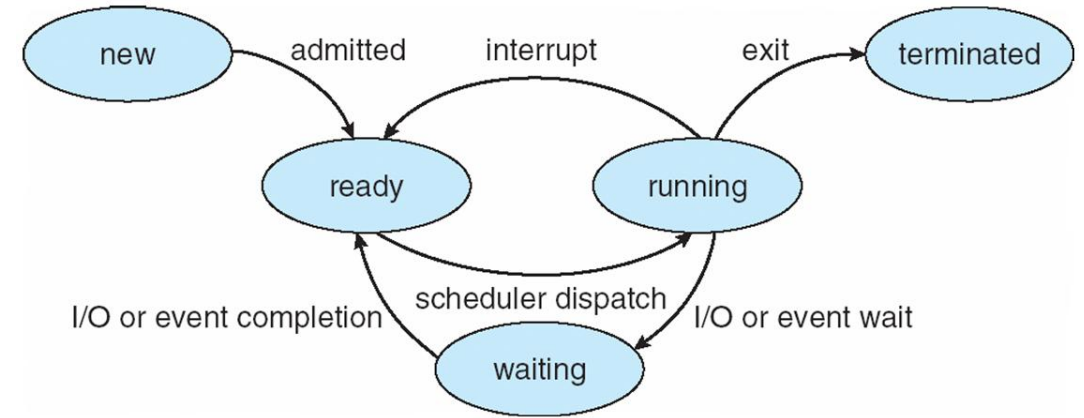
CPU Scheduler

- CPU scheduler (=short term scheduler)
 - CPU scheduler selects a process from ready queue and allocates a CPU core to one of them
- Implementation of ready queue
 - FIFO queue, priority queue, a tree, or simply an ordered linked list
 - Each process is represented by PCB



When Scheduling Occurs?

- 1. Switches from running to waiting state
 - ex) I/O request or invocation of `wait()`
- 2. Switches from running to ready state
 - ex) interrupt
- 3. Switches from waiting to ready
 - ex) completion of I/O
- 4. Terminates



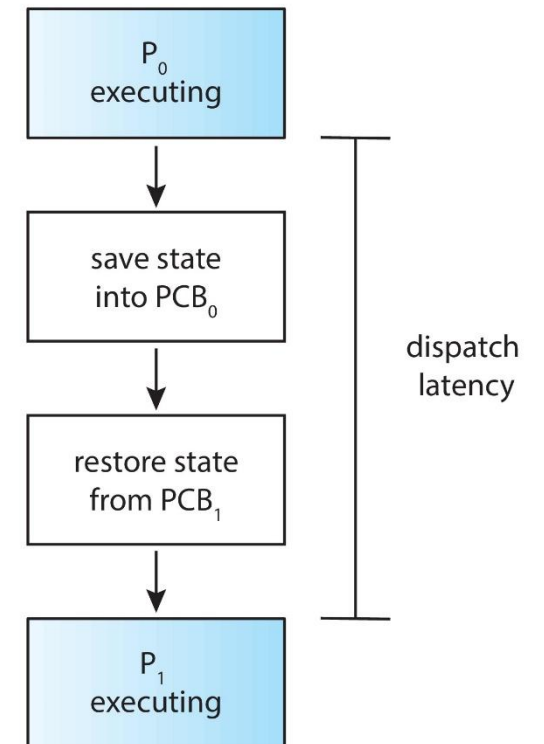
- For situations 1 and 4, there is no choice
- For situations 2 and 3, there is a choice

Preemptive Scheduling

- Non-preemptive (or cooperative) scheduling
 - Scheduling can occur at 1, 4 only
 - Running process is not interrupted
- Preemptive scheduling
 - Scheduling can occurs at 1, 2, 3, 4
 - Scheduling may occur while a process is running
 - ex) interrupt, process with higher priority
 - Requires H/W support and shared data handling
 - Preemptive scheduling can result in **race conditions** when data are shared among several processes

Dispatcher

- Dispatcher: a module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- Dispatch latency: time it takes for the dispatcher to stop one process and start another running



Scheduling Criteria

- **CPU utilization:** keep the CPU as busy as possible
- **Throughput:** # of processes completed per time unit
- **Turnaround time:** interval from submission of a process to its completion
- **Waiting time:** sum of periods spent waiting in ready queue
- **Response time:** time from submission of a request to first response
 - Importance of each criterion vary with systems
- Measure to optimize
 - Average / minimum / maximum value
 - Variance

Agenda

- Basic Concepts
- **Scheduling Algorithms**
- Real-Time CPU Scheduling
- Linux Scheduler



Scheduling Algorithms

- First-come, first-served (FCFS) scheduling
- Shortest-job-first (SJF) scheduling
- Priority scheduling
- Round-robin scheduling
- Multilevel queue scheduling
- Multiple feedback-queue scheduling

First-Come, First-Served Scheduling

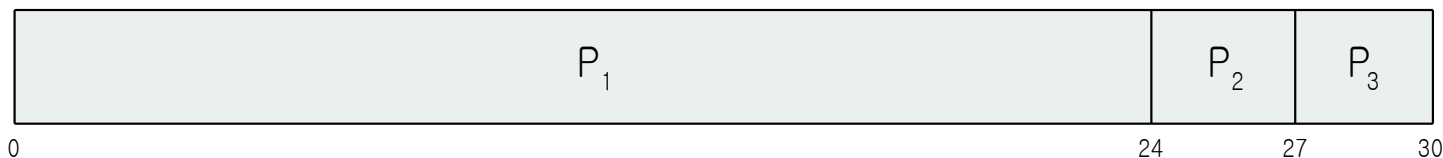
- Process that requests CPU first, is allocated CPU first
 - Non-preemptive scheduling
 - Simplest scheduling method
- Sometimes average waiting time is quite long
 - CPU, I/O utilities are inefficient
 - Ex) Three processes arrived at time 0

Process	Burst Time	Waiting Time
P ₁	24	0
P ₂	3	24
P ₃	3	27

Average waiting time:
 $(0 + 24 + 27) / 3 = 17 \text{ msec}$

What if the processes arrive in the order:
P₂ , P₃ , P₁?

Gantt Chart

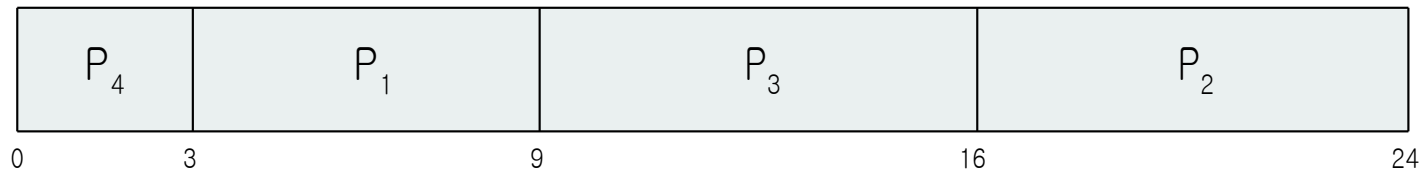


Shortest-Job-First Scheduling

- Assign to the process with the smallest next CPU burst
 - Ex) Four processes arrived at time 0

Process	Burst Time	Waiting Time
P ₁	6	3
P ₂	8	16
P ₃	7	9
P ₄	3	0

Average waiting time:
 $(3 + 16 + 9 + 0) / 4 = 7$ msec



- SJF algorithm is **optimal** in minimum waiting time
- Problem: difficult to know length of next CPU burst

Shortest-Job-First Scheduling

- Predicting next CPU burst from history

- Exponential averaging

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

- t_n : actual length of n-th CPU burst
 - τ_n : prediction for n-th CPU burst
 - α : a coefficient between 0 and 1
 - $\alpha = 0$: recent history has no effect
 - $\alpha = 1$: only recent history matters
 - Usually, $\alpha = 0.5$
 - If we expand the formula, we get:

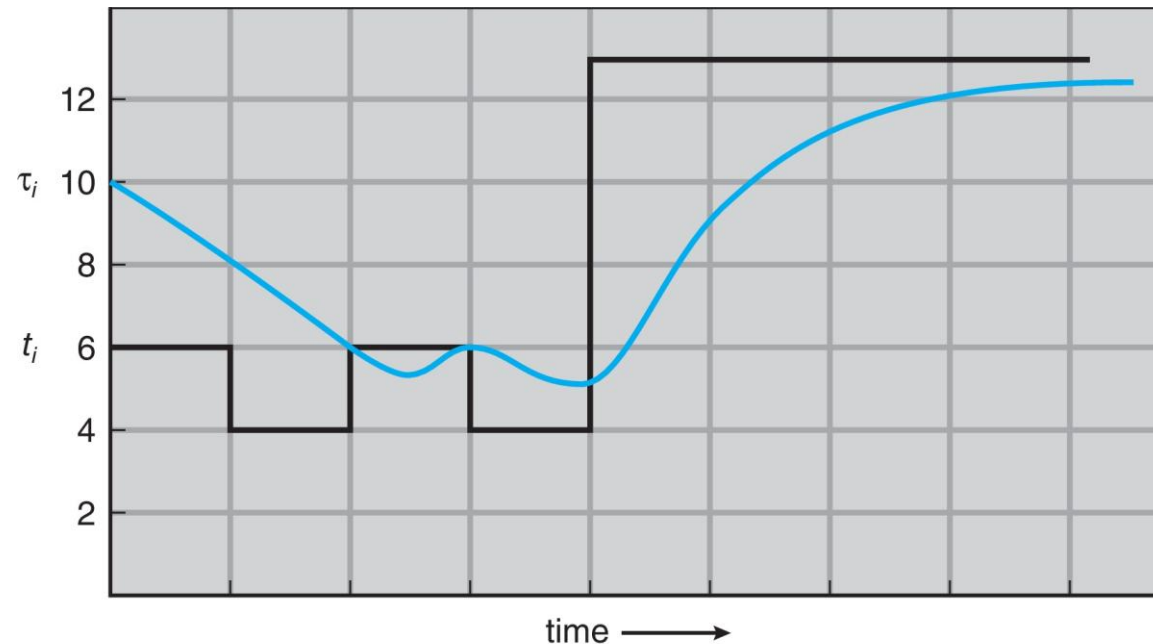
$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \cdots + (1 - \alpha)^j \alpha t_{n-j} + \cdots + (1 - \alpha)^{n+1} \tau_0$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor

Shortest-Job-First Scheduling

- Prediction of CPU Burst

$$\tau_{n+1} = \alpha t_n + (1-\alpha) \tau_n$$



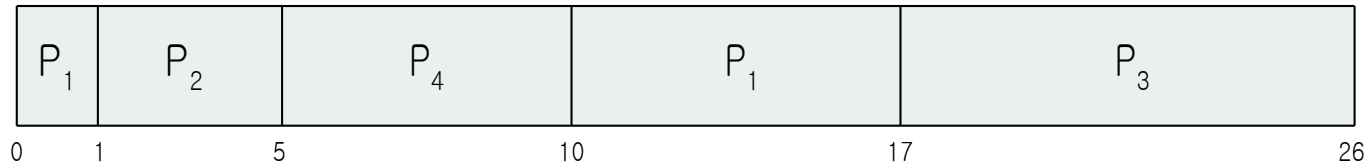
CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

Shortest-Job-First Scheduling

- Preemptive version of SJF scheduling
 - Shortest-remaining-time-first scheduling

Process	Arrival Time	Burst Time	Waiting Time
P ₁	0	8	9 (= 10 - 1)
P ₂	1	4	0 (= 1 - 1)
P ₃	2	9	15 (= 17 - 2)
P ₄	3	5	2 (= 5 - 3)

Average waiting time:
 $(9 + 0 + 15 + 2) / 4 = 26 / 4 = 6.5$



Round-Robin Scheduling

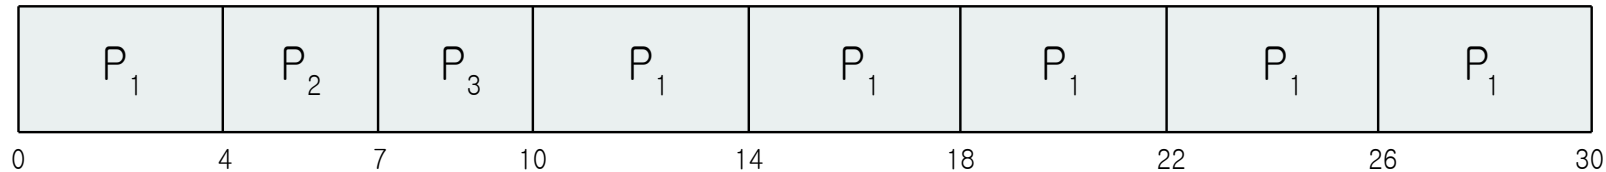
- Similar to FCFS, but it's **preemptive**
 - Designed for time-sharing systems
 - CPU time is divided into **time quantum** (or **time slice**)
 - A time quantum is 10~100 msec.
Cf. switching latency: 10 μ sec
 - Ready queue is treated as **circular queue**
 - CPU scheduler goes around the ready queue and **allocate CPU time up to 1 time quantum**

Round-Robin Scheduling

- Example of RR scheduling with time quantum = 4
 - Ex) Three processes arrived at time 0

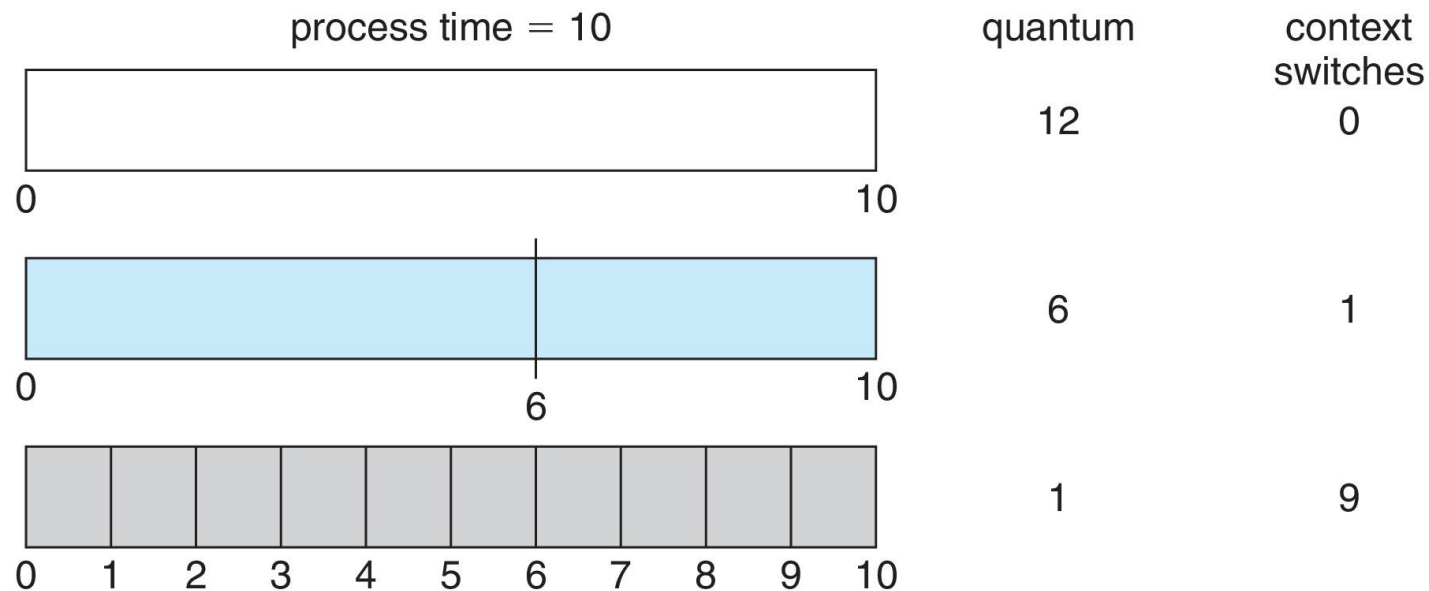
Process	Burst Time	Waiting Time
P ₁	24	6
P ₂	3	4
P ₃	3	7

Average waiting time:
 $(6 + 4 + 7) / 3 = 5.66$



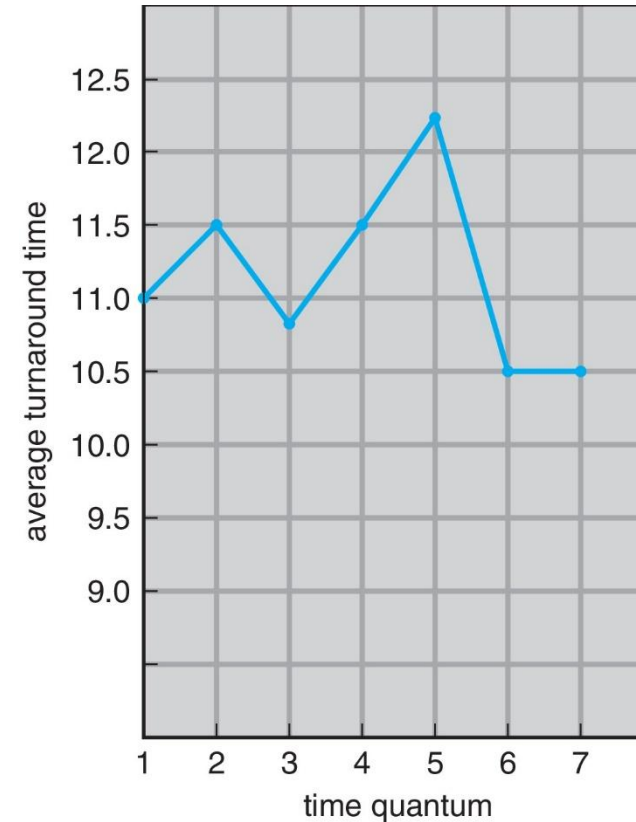
Round-Robin Scheduling

- Performance of RR scheduling heavily depends on size of time quantum
 - Time quantum is small: processor sharing
 - Time quantum is large: FCFS
- Context switching overhead depends on size of time quantum



Round-Robin Scheduling

- Turnaround time also depends on size of time quantum
 - Average turnaround time is not proportional nor inverse-proportional to size of time quantum
 - Average turnaround time is improved if most processes finish their next CPU burst in a single time quantum
 - However, too long time quantum is not desirable
 - A rule of thumb: about 80% of CPU burst should be shorter than time quantum



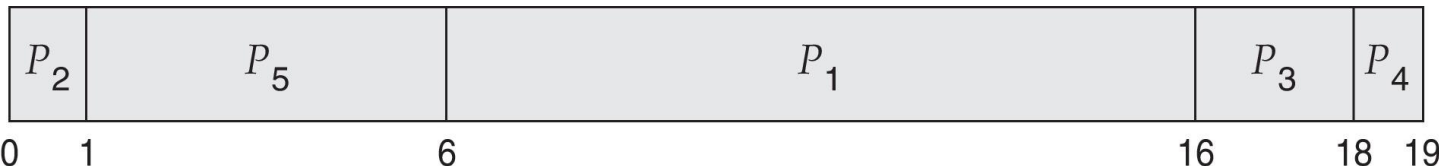
process	time
P_1	6
P_2	3
P_3	1
P_4	7

Priority Scheduling

- CPU is allocated the process with the highest priority
 - Each process has its priority
 - In this text, lower number means higher priority
 - Equal-priority processes: FCFS
 - Ex) Five processes arrived at time 0

Process	Burst Time	Priority	Waiting Time
P ₁	10	3	6
P ₂	1	1	0
P ₃	2	4	16
P ₄	1	5	18
P ₅	5	2	1

Average waiting time:
 $(6 + 0 + 16 + 18 + 1) / 5 = 8.2$



Priority Scheduling

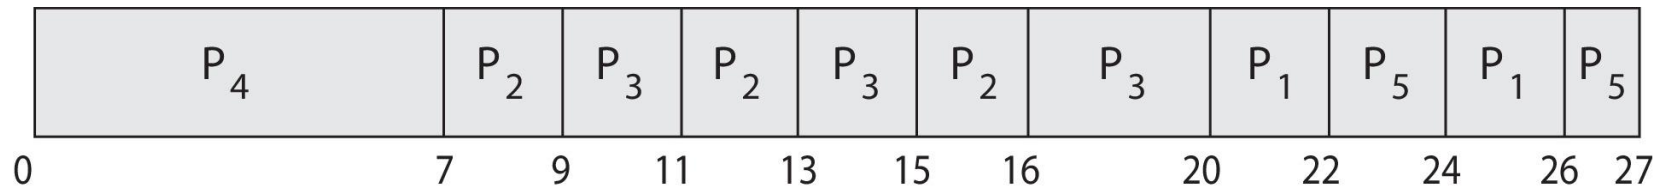
- Priority can be assigned **internally** and **externally**
 - Internally: determined by measurable quantity or qualities
 - Time limit, memory requirement, # of open files, ratio of I/O burst and CPU burst, ...
 - Externally: importance, political factors
- Priority scheduling can be either preemptive or non-preemptive.
- Major problems
 - **Indefinite blocking** (= **starvation**) of processes with lower priorities
 - Solution: **aging** (gradually increase priority of processes waiting for long time)

Priority Scheduling

- Priority scheduling with Round-Robin scheduling
 - Run the process with the highest priority
 - Processes with the same priority run round-robin
 - Ex) 2 msec time quantum for RR

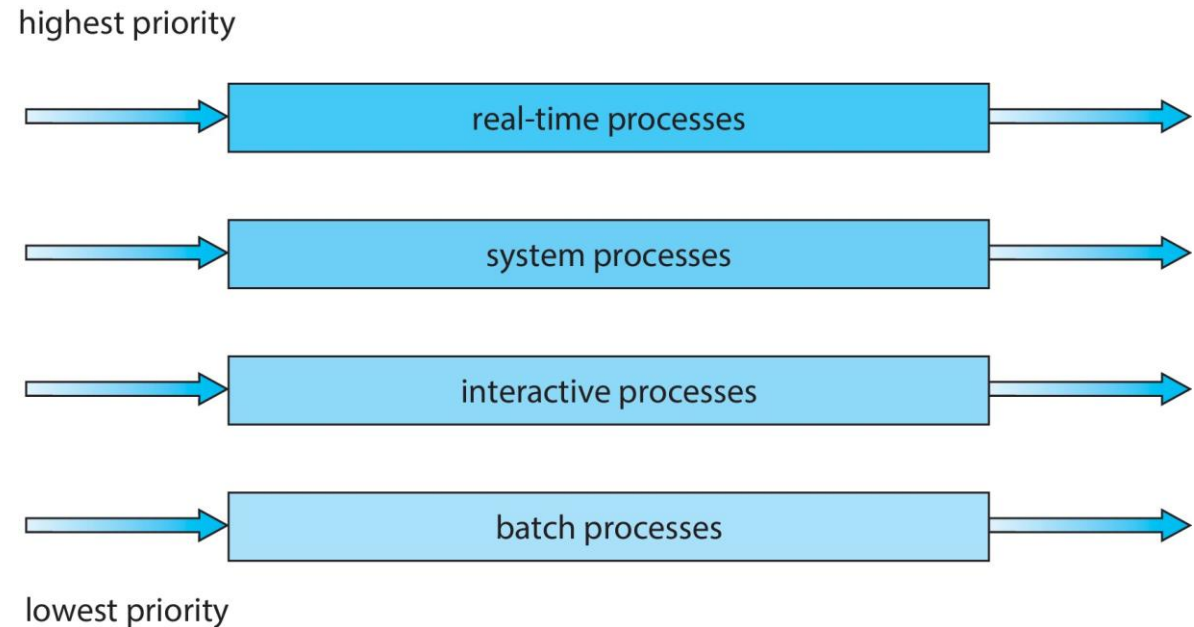
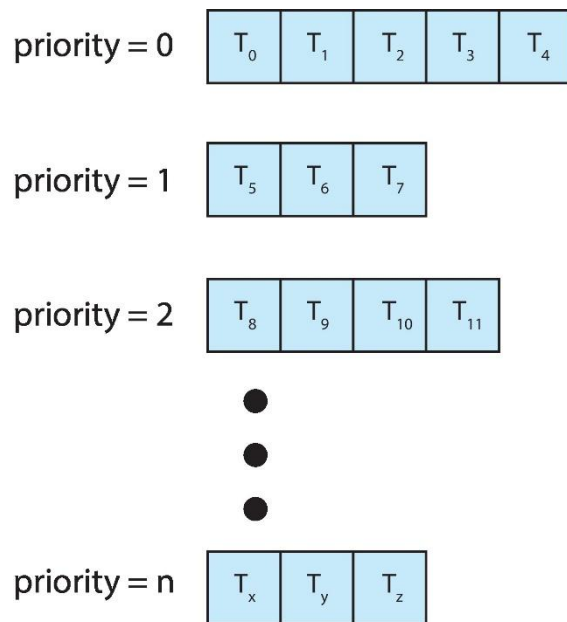
Process	Burst Time	Priority	Waiting Time
P ₁	4	3	22 (=20+2)
P ₂	5	2	11 (=7+2+2)
P ₃	8	2	12 (=9+2+1)
P ₄	7	1	0
P ₅	3	3	24 (=22+2)

Average waiting time:
 $(22 + 11 + 12 + 0 + 24) / 5 = 13.8$



Multilevel Queue Scheduling

- Classify processes into different groups and apply different scheduling
 - Memory requirement, priority, process type, ...
- Partition ready queue into several separate queues



Multilevel Queue Scheduling

- Each queue has its own scheduling algorithm
- Scheduling among queues
 - Fixed-priority preemptive scheduling
 - A process in lower priority queue can run only when all of higher priority queues all empty
 - Time-slice among queues

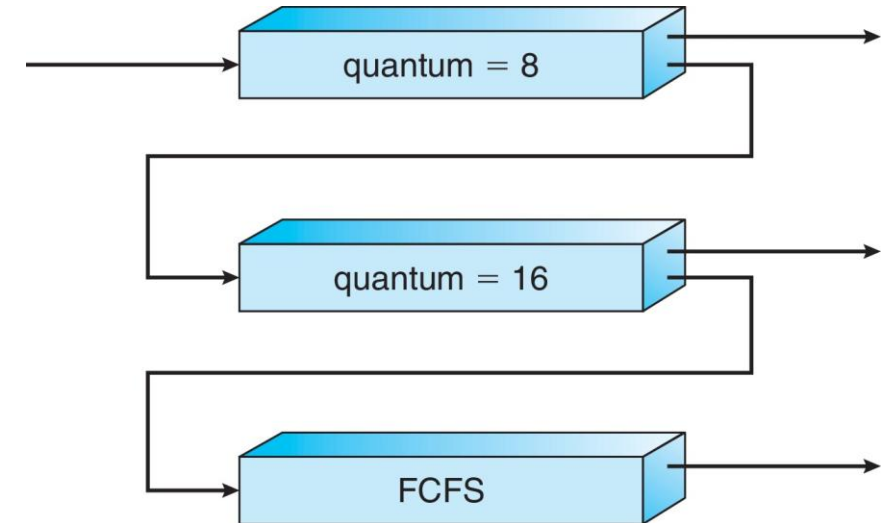
Ex) foreground queue (interactive processes): 80% of the CPU time for RR
background queue (batch processes): 20% of the CPU time for FCFS

Multilevel Feedback-Queue Scheduling

- Similar to multilevel queue scheduling, but a process can move between queues
- Idea: separate processes according to characteristics of their CPU bursts.
 - If a process uses too much CPU time, move it to lower priority queue
 - I/O-bound, interactive processes are in higher priority queues

Multilevel Feedback-Queue Scheduling

- Example of multilevel feedback queue
 - Ready queue consists of three queues (0~2)
 - Q_0 : RR with time quantum 8 msec
 - Q_1 : RR with time quantum 16 msec
 - Q_2 : FCFS
- A new process is put in Q_0 .
If it exceeds time limit, it moves to lower priority queue



Multilevel Feedback-Queue Scheduling

- Parameters to define a multilevel feedback-queue scheduler
 - # of queues
 - Scheduling algorithm for each queue
 - Method to determine when to upgrade a process to higher priority queue
 - Method to determine when to demote a process to lower priority queue
 - Method to determine which queue a process will enter when it needs service
- The most complex algorithm

Agenda

- Basic Concepts
- Scheduling Algorithms
- **Real-Time CPU Scheduling**
- Operating System Example



Real-time CPU Scheduling

- Real-time Operating Systems (RTOS)

- OS intended to serve real-time systems

Ex) Antilock brake system (ABS) requires latency of 3~5 msec

- Soft real-time systems

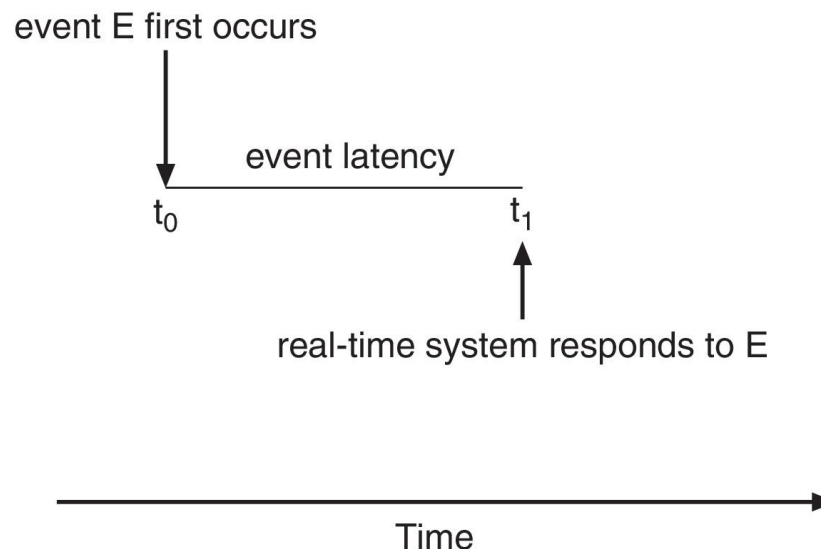
- No guarantee as to when a critical real-time process will be scheduled.
 - Guarantee only that the process will be given preference over noncritical processes.

- Hard real-time systems

- A task must be serviced by its deadline
 - Service after the deadline has expired is the same as no service at all.

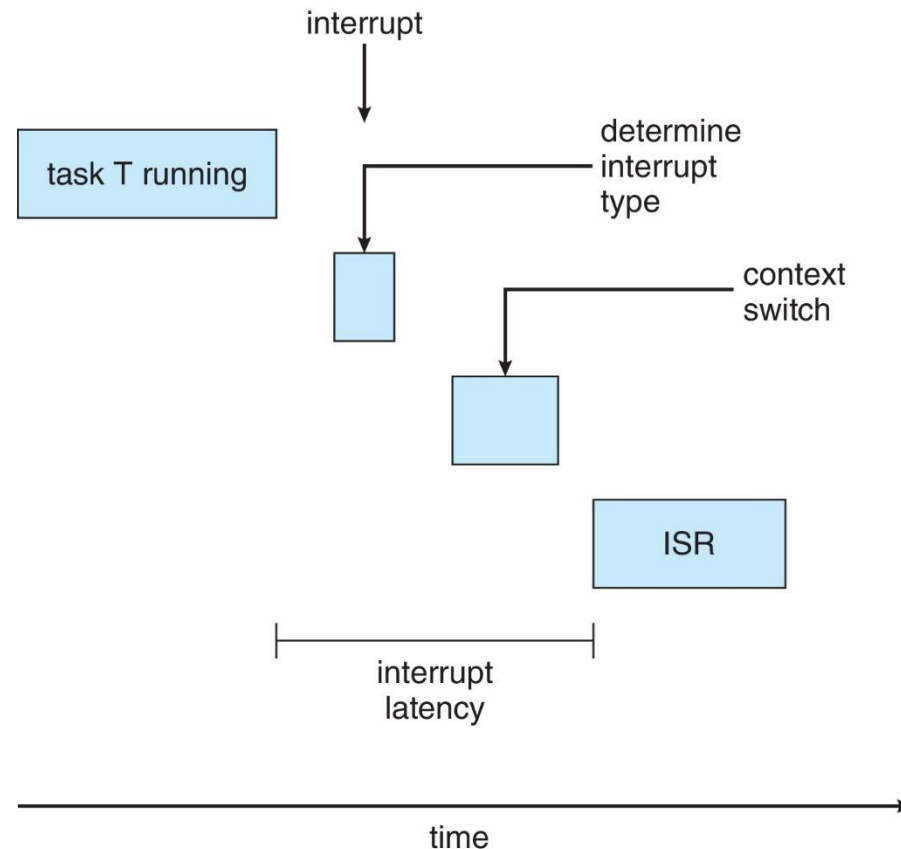
Minimizing Latency

- Most real-time systems are **event-driven system**
 - When an event occurs, the system must respond to and service it as quickly as possible
- **Event latency**: the amount of time that elapses from when an event occurs to when it is serviced
 - Interrupt latency, dispatch latency



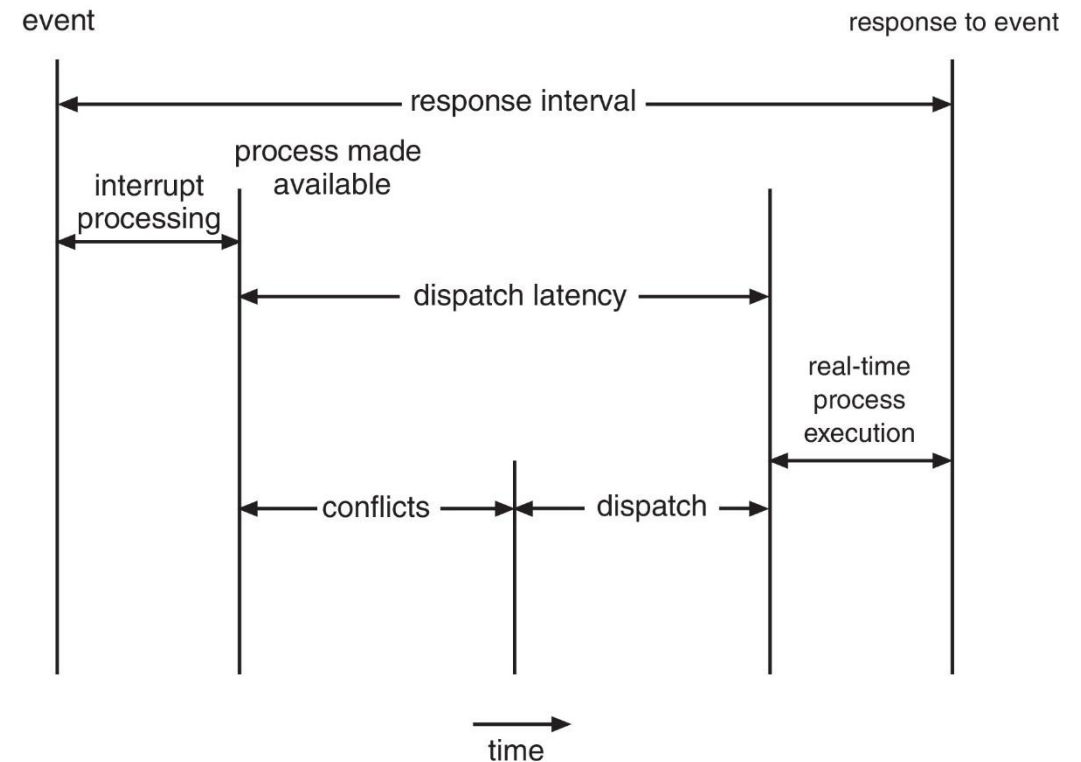
Interrupt Latency

- **Interrupt latency**: the period of time from the arrival of interrupt at the CPU to the start of the routine that services the interrupt.



Dispatch Latency

- **Dispatch latency**: the amount of time required for the scheduling dispatcher to stop one process and start another
 - Preemptive kernels are the most effective technique to keep dispatch latency low
- Conflict phase of dispatch latency
 - 1. Preemption of any process running in kernel mode
 - 2. Release by low-priority process of resources needed by high-priority processes



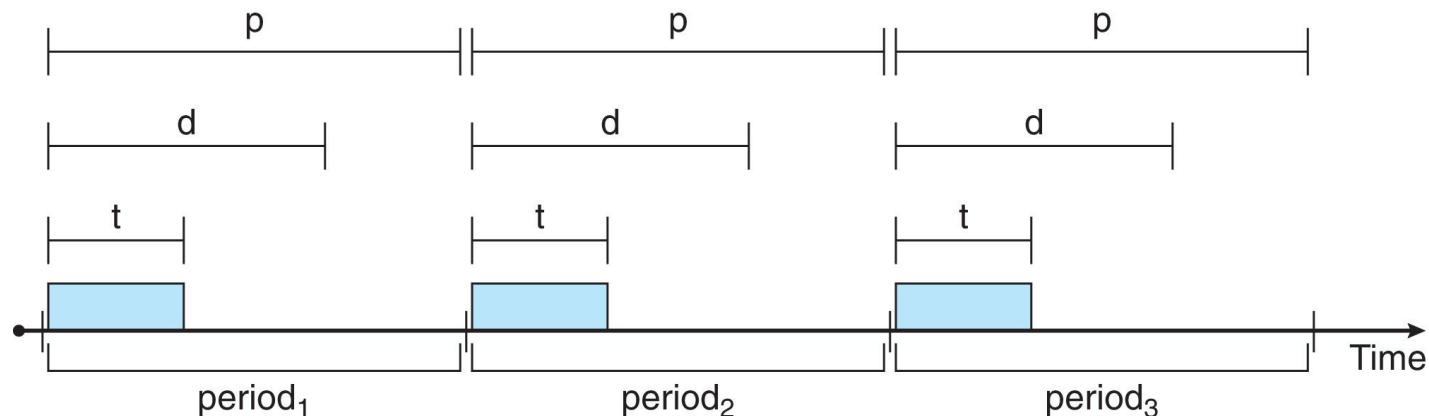
Priority-based Scheduling

- The most important feature of RTOS is to respond immediately to a real-time process
 - The scheduler should support priority-based preemptive algorithm
 - Most OS assign highest priority to real-time processes
 - Preemptive, priority-based scheduler only guarantees soft real-time functionality.
- Hard real-time systems must further guarantee that real-time tasks will be serviced in accord with their deadline requirements
 - Requires additional scheduling features.

Periodic Process

- Periodic processes require the CPU at constant intervals (periods).
 - Processing time: t
 - Deadline: d (by which it must be serviced by the CPU)
 - Period: p
 - Rate of a periodic task: $1/p$

$$0 \leq t \leq d \leq p$$



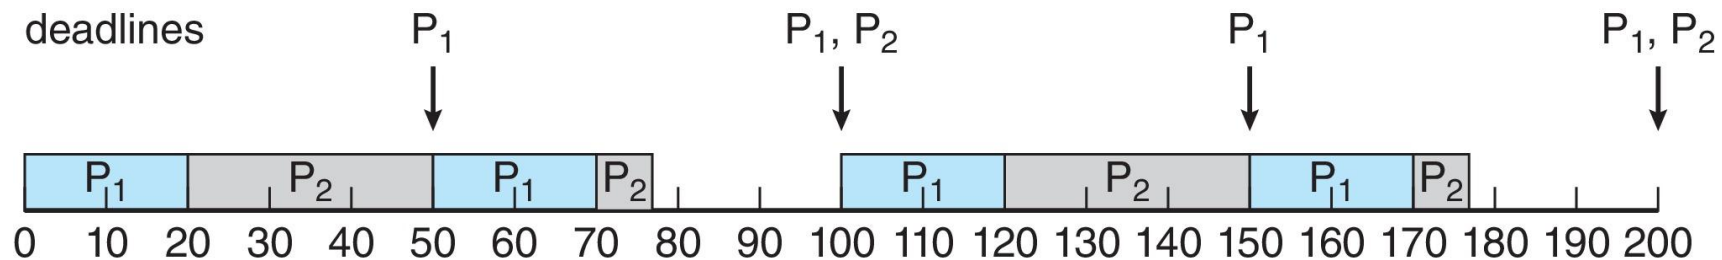
Scheduling with Deadline Requirement

- A process may have to announce its deadline requirements to the scheduler
- Using an admission-control algorithm, the scheduler does one of two things:
 - **Admits the process**, guaranteeing that the process will complete on time, or
 - **Rejects the request** as impossible if it cannot guarantee that the task will be serviced by its deadline.

Rate-Monotonic Scheduling

- Rate-monotonic scheduling algorithm
 - Schedules periodic tasks using **a static priority policy with preemption**
 - Each periodic task is assigned a priority **inversely based on its period**
 - Shorter periods \rightarrow high priority
 - Longer periods \rightarrow lower priority
 - The processing time of a periodic process is assumed the same for each CPU burst

Process $P_1: p_1 = 50, t_1 = 20$, Process $P_2: p_2 = 100, t_2 = 35$

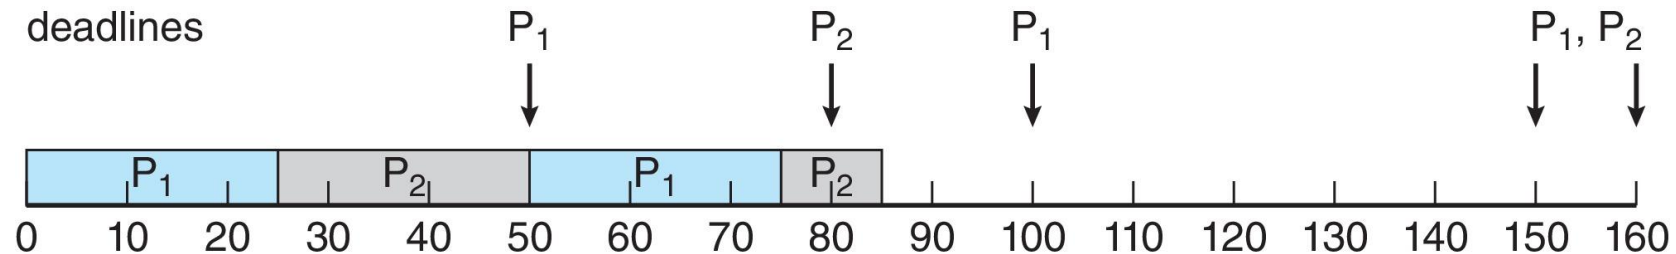


Rate-Monotonic Scheduling

- Rate-monotonic scheduling is considered optimal
 - If a set of processes cannot be scheduled by this algorithm, it cannot be scheduled by any other algorithm that assigns static priorities.
- Limitation
 - CPU utilization is bounded, and it is not always possible to maximize CPU resources fully.
 - Worst-case CPU utilization: $N(2^{1/N} - 1)$
 - N : # of processes

Process P_1 : $p_1 = 50$, $t_1 = 25$, Process P_2 : $p_2 = 80$, $t_2 = 35$

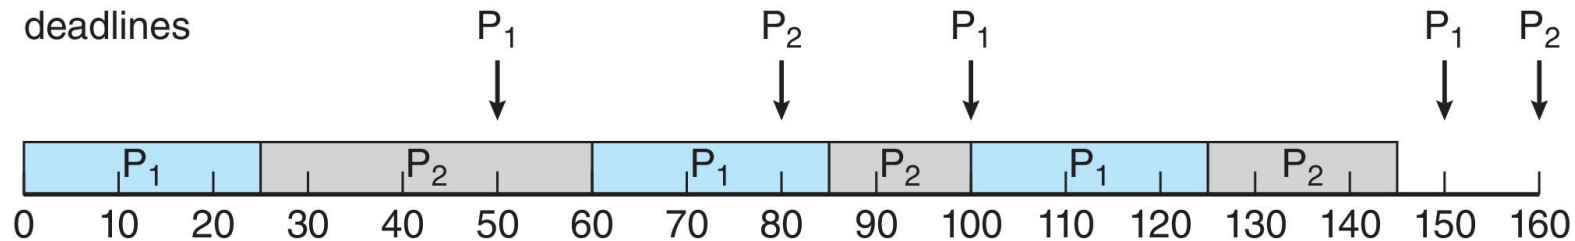
Process P2 misses
finishing its deadline
at time 80



Earliest-Deadline-First Scheduling

- Earliest-deadline-first (EDF) scheduling
 - Priorities are assigned according to deadline
 - The earlier the deadline, the higher the priority
 - The later the deadline, the lower priority

Process P_1 : $p_1 = 50$, $t_1 = 25$, Process P_2 : $p_2 = 80$, $t_2 = 35$



Earliest-Deadline-First Scheduling

- Requirements

- Does not require that processes be periodic, nor must a process require a constant amount of CPU time per burst
- When a process becomes runnable, it must announce its deadline requirements to the system

- Theoretically optimal

- Theoretically, it can schedule processes so that each process can meet its deadline requirements and CPU utilization will be 100 percent.

Proportional Share Scheduling

- Proportional share schedulers
 - Allocates T shares among all applications.
 - An application can receive N shares of time, thus ensuring that the application will have N / T of the total processor time
 - Must work in conjunction with an admission-control policy to guarantee that an application receives its allocated shares of time

Agenda

- Basic Concepts
- Scheduling Algorithms
- Real-Time CPU Scheduling
- Linux Scheduler



Linux Scheduler

■ History

- Traditional UNIX scheduling algorithm (ver. < 2.5)
 - Not adequate support for SMP system
 - Not scale well as # of tasks increases
- O(1) scheduling algorithm (ver. \geq 2.5)
 - O(1) complexity
 - Increased support for SMP
 - Poor response time for the interactive processes
- Completely Fair Scheduler (CFS) (ver. \geq 2.6)

Linux Scheduler

- Based on scheduling classes
 - Each class is assigned a specific priority
 - Default scheduling class (CFS algorithm)
 - Static priority 100 ~ 139
 - Real-time scheduling class (SCHED_FIFO, SCHED_RR)
 - Static priority 0 ~ 99
 - If necessary, new scheduling classes can be added
 - Allows different scheduling algorithms based on the needs of the system (e.g. server vs. mobile)
 - Scheduler selects highest-priority task belonging to highest-priority class

Linux CFS Scheduler

- CFS scheduler assigns a proportion of CPU processing time to each task
 - The portion is calculated from nice value
 - Nice value: relative priority in $[-20, +19]$, default is 0
 - Higher value represents lower priority ('nice' to other tasks)
 - Proportions of CPU time are allocated from the value of targeted latency
 - Targeted latency: interval of time during which every runnable task should run at least once
 - Can increase if the number of active tasks in the system grows beyond a certain threshold

Linux CFS Scheduler

- Records how long each task has run
 - per task variable **vruntime** (virtual run time)
 - Normal priority tasks: $\text{vruntime} = \langle \text{actual physical run time} \rangle$
- vruntime of each task is associated with a **decay factor**
 - Lower priority task: higher decay factor
 - Higher priority task: lower decay factor
 - ➔ Higher priority tasks are likely to have smaller vruntime
- The scheduler simply selects the task that has **the smallest vruntime value**
- Higher-priority task can preempt a lower-priority task.

Linux CFS Scheduler

- I/O-bound tasks: run only for short periods before blocking for additional I/O
 - vruntime will be lower
 - CPU-bound tasks: exhaust its time period whenever it has an opportunity to run on a processor
 - vruntime will be higher
- Gives I/O-bound tasks higher priority than CPU-bound tasks